Contents

List of Figures

List of Acronyms

Adam	Adaptive Moment Estimation (an optimisation algorithm)
AI	Artificial Intelligence
Aleph	A Learning Engine for Proposing Hypotheses (an ILP system)
BotGNN	Bottom-Graph Neural Network
BK	Background Knowledge
CNN	Convolutional Neural Network
CONV	Convolution (used for a block or a layer)
DL	Deep Learning
DNN	Deep Neural Network
DRM	Deep Relational Machine
GNN	Graph Neural Network
ILP	Inductive Logic Programming
MDIE	Mode-Directed Inverse Entailment
ML	Machine Learning
MLP	Multilayer Perceptron
NCI	National Cancer Institute
NN	Neural Network
POOL	Pooling (used for a block or a layer)
RNN	Recurrent Neural Network
SGD	Stochastic Gradient Descent
SMILES	Simplified Molecular-Input Line-Entry System
VAE	Variational Autoencoder
VEGNN	Vertex-Enriched Graph Neural Network
XAI	Explainable Artificial Intelligence